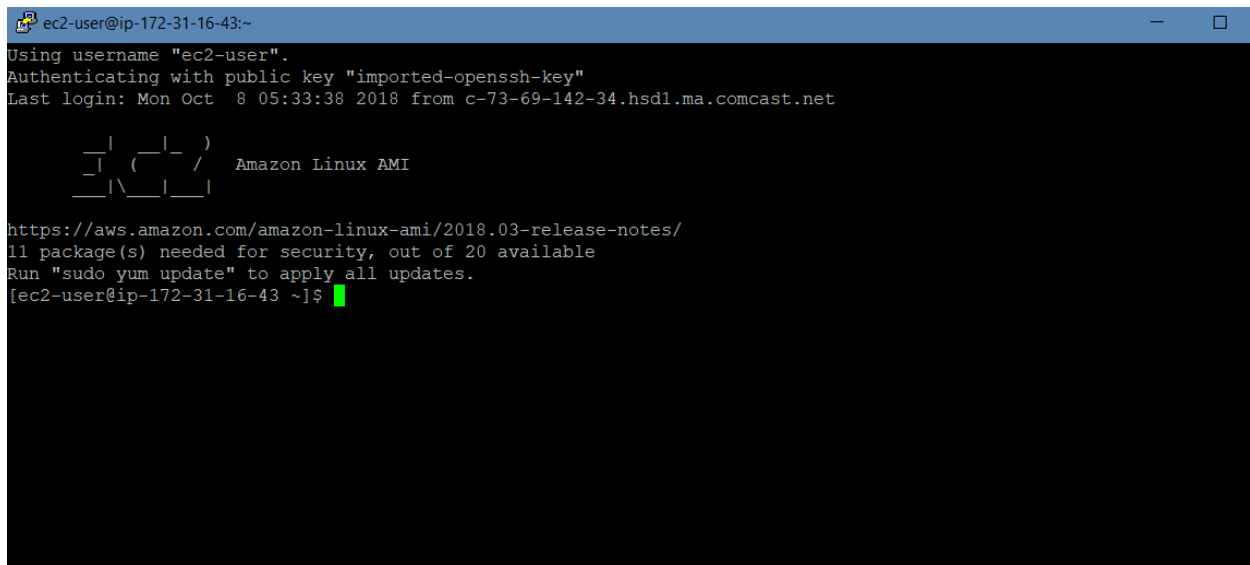


Installing Kubernetes on AWS EC2

Prerequisites:

- Ubuntu instance (You may use other linux instance as well)
- AWS-cli setup
- S3 bucket



```
ec2-user@ip-172-31-16-43:~  
Using username "ec2-user".  
Authenticating with public key "imported-openssh-key"  
Last login: Mon Oct  8 05:33:38 2018 from c-73-69-142-34.hsd1.ma.comcast.net  
  
  _   _   _  
 _(_)_/   ) Amazon Linux AMI  
/_/_/_/___/_____  
  
https://aws.amazon.com/amazon-linux-ami/2018.03-release-notes/  
11 package(s) needed for security, out of 20 available  
Run "sudo yum update" to apply all updates.  
[ec2-user@ip-172-31-16-43 ~]$
```

Install kubectl

On ubuntu instance, make sure you have AWS cli and KOPS setup. We shall also need **kubectl** (Kubernetes cli)

- Install Kubectl on Linux:

```
curl -LO https://storage.googleapis.com/kubernetes-release/release/$(curl -s  
https://storage.googleapis.com/kubernetes-  
release/release/stable.txt)/bin/linux/amd64/kubectl  
chmod +x ./kubectl  
sudo mv ./kubectl /usr/local/bin/kubectl
```

```
[ec2-user@ip-172-31-16-43 ~]$ curl -LO https://storage.googleapis.com/kubernetes-release/release/$(curl -s https://storage.googleapis.com/kubernetes-release/release/stable.txt)/bin/linux/amd64/kubectl
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
100 54.6M  100 54.6M    0     0  49.4M    0  0:00:01  0:00:01 --:--:-- 49.4M
[ec2-user@ip-172-31-16-43 ~]$ chmod +x ./kubectl
[ec2-user@ip-172-31-16-43 ~]$ sudo mv ./kubectl /usr/local/bin/kubectl
[ec2-user@ip-172-31-16-43 ~]$
```

Now, let's install kops on ubuntu box:

wget <https://github.com/kubernetes/kops/releases/download/1.6.1/kops-linux-amd64>

chmod +x kops-linux-amd64

sudo mv kops-linux-amd64 /usr/local/bin/kops

```
[ec2-user@ip-172-31-16-43 ~]$ wget https://github.com/kubernetes/kops/releases/download/1.6.1/kops-linux-amd64
--2018-10-08 16:56:06-- https://github.com/kubernetes/kops/releases/download/1.6.1/kops-linux-amd64
Resolving github.com (github.com)... 192.30.253.113, 192.30.253.112
Connecting to github.com (github.com)|192.30.253.113|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://github-production-release-asset-2e65be.s3.amazonaws.com/62091339/4e48c984-4eca-11e7-9a28-7434a10577bd?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAIWNJYAX4CSVEH53A%2F20181008%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20181008T165606Z&X-Amz-Expires=300&X-Amz-Signature=ce73082057315050d3740737fc2f1079bd8d68ab1bd7a4caf6a4e21c832989aa&X-Amz-SignedHeaders=host&actor_id=0&response-content-disposition=attachment%3B%20filename%3Dkops-linux-amd64&response-content-type=application%2Foctet-stream [following]
--2018-10-08 16:56:06-- https://github-production-release-asset-2e65be.s3.amazonaws.com/62091339/4e48c984-4eca-11e7-9a28-7434a10577bd?X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAIWNJYAX4CSVEH53A%2F20181008%2Fus-east-1%2Fs3%2Faws4_request&X-Amz-Date=20181008T165606Z&X-Amz-Expires=300&X-Amz-Signature=ce73082057315050d3740737fc2f1079bd8d68ab1bd7a4caf6a4e21c832989aa&X-Amz-SignedHeaders=host&actor_id=0&response-content-disposition=attachment%3B%20filename%3Dkops-linux-amd64&response-content-type=application%2Foctet-stream
Resolving github-production-release-asset-2e65be.s3.amazonaws.com (github-production-release-asset-2e65be.s3.amazonaws.com)... 52.216.81.144
Connecting to github-production-release-asset-2e65be.s3.amazonaws.com (github-production-release-asset-2e65be.s3.amazonaws.com)|52.216.81.144|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 72731008 (69M) [application/octet-stream]
Saving to: 'kops-linux-amd64'

kops-linux-amd64      100%[=====>] 69.36M  69.4MB/s  in 1.0s

2018-10-08 16:56:07 (69.4 MB/s) - 'kops-linux-amd64' saved [72731008/72731008]

[ec2-user@ip-172-31-16-43 ~]$ chmod +x kops-linux-amd64
[ec2-user@ip-172-31-16-43 ~]$ sudo mv kops-linux-amd64 /usr/local/bin/kops
[ec2-user@ip-172-31-16-43 ~]$
```

Create Route53 domain for the cluster

kubernetes make use of DNS for discovery within the cluster so that you can reach out kubernetes-API-server from clients.

Create a hosted zone on Route53, say, ***k8s.appychip.vpc***. The API server endpoint will then be ***api.k8s.appychip.vpc***

Create a S3 bucket

Now, create a S3 bucket to store the configuration for the cluster. Make sure the instance have right role to access S3 and Route53:

```
$ aws s3 mb s3://clusters.k8s.appychip.vpc
```

```
[ec2-user@ip-172-31-16-43 ~]$ aws s3 mb s3://clusters.k8s.aws.amazon.com
make_bucket: clusters.k8s.aws.amazon.com
[ec2-user@ip-172-31-16-43 ~]$
```

Expose environment variable:

```
$ export KOPS_STATE_STORE=s3://clusters.k8s.appychip.vpc
```

Create Kubernetes Cluster

Now comes the interesting part to create the cluster. You can reuse existing VPC (kops will create a new subnet in this VPC) by providing the **vpc-id** option. The following command will give you details what all things are going to happen:

```
$ kops create cluster --cloud=aws --zones=us-east-1d --
name=useast1.k8s.appychip.vpc --dns-zone=appychip.vpc --dns private
```

```

[ec2-user@ip-172-31-16-43 ~]$ aws s3 mb s3://clusters.k8s.aws.amazon.com
make_bucket: clusters.k8s.aws.amazon.com
[ec2-user@ip-172-31-16-43 ~]$ export KOPS_STATE_STORE=s3://clusters.k8s.aws.amazon.com
[ec2-user@ip-172-31-16-43 ~]$ kops create cluster --cloud=aws --zones=us-east-2b --name=useast2b.k8s.aws.amazon.com --
dns-zone=amazon.aws.com --dns private
I1008 17:01:17.484529    2689 create_cluster.go:833] Using SSH public key: /home/ec2-user/.ssh/id_rsa.pub
I1008 17:01:17.640318    2689 subnets.go:183] Assigned CIDR 172.20.32.0/19 to subnet us-east-2b
W1008 17:01:17.731079    2689 populate_instancegroup_spec.go:209] "m3.medium" instance is not available in region "us-
east-2", will set master to "c4.large" instead
Previewing changes that will be made:

*****

A new kops version is available: 1.8.1

This version of kops is no longer supported; upgrading is required
(you can bypass this check by exporting KOPS_RUN_OBSOLETE_VERSION)

More information: https://github.com/kubernetes/kops/blob/master/permalinks/upgrade_kops.md#1.8.1

*****

kops upgrade is required
[ec2-user@ip-172-31-16-43 ~]$ █

```

UPGRADE KOPS

wget -O kops [https://github.com/kubernetes/kops/releases/download/\\$\(curl -s https://api.github.com/repos/kubernetes/kops/releases/latest | grep tag_name | cut -d '"' -f 4\)/kops-linux-amd64](https://github.com/kubernetes/kops/releases/download/$(curl -s https://api.github.com/repos/kubernetes/kops/releases/latest | grep tag_name | cut -d ')

NOTE: Make sure you have ssh keys already generated otherwise it will throw an error.

Generating SSH key

```
[ec2-user@ip-172-31-16-43 ~]$ ssh-keygen
Generating public/private rsa key pair.
Enter file in which to save the key (/home/ec2-user/.ssh/id_rsa):
/home/ec2-user/.ssh/id_rsa already exists.
Overwrite (y/n)? Y
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/ec2-user/.ssh/id_rsa.
Your public key has been saved in /home/ec2-user/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:vgBlnIxdqHguB0UYlGYeeHZ++PE2AnmqF8jBnPhLHps ec2-user@ip-172-31-16-43
The key's randomart image is:
+---[RSA 2048]-----+
|  oo=.  ..          |
| . O o=.o          |
| X Bo+*           |
| . O.B.+          |
| o.* *So          |
| B.=.o +          |
| o O...o .        |
| E .. .          |
| . . .          |
+---[SHA256]-----+
[ec2-user@ip-172-31-16-43 ~]$
```

Now to actually create cluster run the following command:

```
kops update cluster useast1.k8s.appychip.vpc --yes
```

Now after the update cluster command has been executed the clusters will be created in AWS EC2

- * validate cluster: kops validate cluster
- * list nodes: kubectl get nodes --show-labels
- * ssh to the master: ssh -i ~/.ssh/id_rsa admin@api.useast1.k8s.appychip.vpc

aws

Services

Resource Groups

+

ADS-Team5

Ohio

Support

EC2 Dashboard

Events

Tags

Reports

Limits

INSTANCES

Instances

Launch Templates

Spot Requests

Reserved Instances

Dedicated Hosts

IMAGES

AMIs

Bundle Tasks

ELASTIC BLOCK STORE

Volumes

Snapshots

Launch Instance

Connect

Actions

Filter by tags and attributes or search by keyword

1 to 4 of 4

	Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS (IPv4)
<input type="checkbox"/>	nodes-useast1-k8s-team5-ads1	i-00cc04035c9ec15aa	t2.medium	us-east-2b	running	2/2 checks ...	None	ec2-18-191-204-163
<input type="checkbox"/>		i-07611338bc1188dc	t2.micro	us-east-2b	running	2/2 checks ...	None	ec2-18-188-247-176
<input type="checkbox"/>	nodes-useast1-k8s-team5-ads1	i-0d0eb9cca9169c8a8	t2.medium	us-east-2b	running	2/2 checks ...	None	ec2-18-221-74.23.us
<input type="checkbox"/>	master-us-east-2b-masters-useast1-k8s-team5-ads1	i-0e9ba85343ec08ec5	c4.large	us-east-2b	running	2/2 checks ...	None	ec2-18-224-171-114

Select an instance above

```
ubuntu@ip-172-31-77-124:~$ kops validate cluster
Using cluster from kubectl context: useast1.k8s.appychip.vpc

Validating cluster useast1.k8s.appychip.vpc

INSTANCE GROUPS
NAME                                ROLE    MACHINETYPE  MIN  MAX  SUBNETS
master-us-east-1d                  Master  m3.medium    1    1    us-east-1d
nodes                              Node    t2.medium    2    2    us-east-1d

NODE STATUS
NAME                                ROLE    READY
ip-172-20-53-62.ec2.internal        master  True
ip-172-20-61-83.ec2.internal        node    False
```

To get nodes

Run—kubectl get nodes

```
ubuntu@ip-172-31-77-124:~$ kubectl get nodes
NAME                                STATUS    AGE       VERSION
ip-172-20-53-62.ec2.internal        Ready     2m        v1.6.2
ip-172-20-55-195.ec2.internal       Ready     53s       v1.6.2
ip-172-20-61-83.ec2.internal        Ready     1m        v1.6.2
ubuntu@ip-172-31-77-124:~$
ubuntu@ip-172-31-77-124:~$
ubuntu@ip-172-31-77-124:~$ kubectl create -f https://rawgit.com/kubernetes/dashboard/master/src/deploy/kubernetes-dashboard.yaml
serviceaccount "kubernetes-dashboard" created
clusterrolebinding "kubernetes-dashboard" created
deployment "kubernetes-dashboard" created
service "kubernetes-dashboard" created
ubuntu@ip-172-31-77-124:~$
ubuntu@ip-172-31-77-124:~$
ubuntu@ip-172-31-77-124:~$
```

Deploying Nginx Container

Let's deploy a simple service made up of some nginx containers:

Create an nginx deployment:

```
$ kubectl run sample-nginx --image=nginx --replicas=2 --port=80
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
sample-nginx-379829228-xb9y3	1/1	Running	0	10s
sample-nginx-379829228-yhd25	1/1	Running	0	10s

```
$ kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
sample-nginx	2	2	2	2	29s

Expose the deployment as service. This will create an ELB in front of those 2 containers and allow us to publicly access them:


```
$ kubectl expose deployment sample-nginx --port=80 --type=LoadBalancer
```

```
$ kubectl get services -o wide
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)
kubernetes	100.64.0.1	<none>	443/TCP
sample-nginx	100.70.129.69	adca6650a60e611e7a66612ae64874d4-175711331.us-east-1.elb.amazonaws.com/	80/TCP

19m run=sample-nginx

There is an ELB running on <http://adca6650a60e611e7a66612ae64874d4-175711331.us-east-1.elb.amazonaws.com> with our nginx containers behind it:



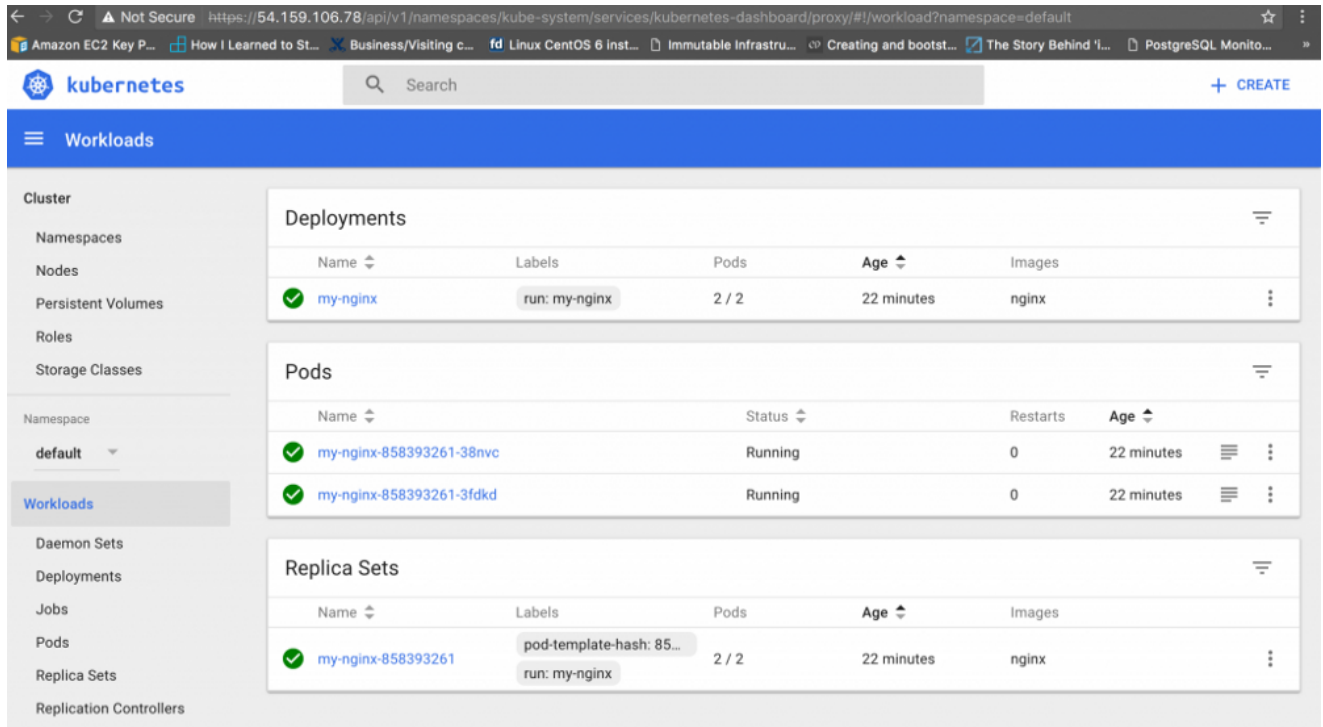
Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

You can also view the UI by accessing master node. Hit master node's IP/Domain in browser, it will ask for credentials. Run command ***kubectl config view*** to see the credentials.



The screenshot shows the Kubernetes dashboard interface. The top navigation bar includes the Kubernetes logo, a search bar, and a '+ CREATE' button. The left sidebar contains a menu with options: Cluster, Namespaces, Nodes, Persistent Volumes, Roles, Storage Classes, Namespace (set to 'default'), Workloads (selected), Daemon Sets, Deployments, Jobs, Pods, Replica Sets, and Replication Controllers. The main content area displays three sections: Deployments, Pods, and Replica Sets. The Deployments section shows a single deployment named 'my-nginx' with 2 pods, labeled 'run: my-nginx', and an age of 22 minutes. The Pods section shows two running pods: 'my-nginx-858393261-38nvc' and 'my-nginx-858393261-3fdkd', both with 0 restarts and an age of 22 minutes. The Replica Sets section shows a single replica set named 'my-nginx-858393261' with 2 pods, labeled 'pod-template-hash: 85...' and 'run: my-nginx', and an age of 22 minutes.

Name	Labels	Pods	Age	Images
my-nginx	run: my-nginx	2 / 2	22 minutes	nginx

Name	Status	Restarts	Age
my-nginx-858393261-38nvc	Running	0	22 minutes
my-nginx-858393261-3fdkd	Running	0	22 minutes

Name	Labels	Pods	Age	Images
my-nginx-858393261	pod-template-hash: 85... run: my-nginx	2 / 2	22 minutes	nginx

To delete the cluster and remove all AWS resources with, run the following command:

```
$ kops delete cluster --name=useast1.k8s.appychip.vpc --yes
```

That's All is required to setup kubernetes on AWS EC2.